

# Customer Churn Prediction

In [1]:

```
import pandas as pd
from matplotlib import pyplot as plt
import numpy as np
%matplotlib inline
```


In [2]:

```
df = pd.read_csv("WA_Fn-UseC_-Telco-Customer-Churn.csv")
df.sample(5)
```

Out[2]:

|      | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLi |
|------|------------|--------|---------------|---------|------------|--------|--------------|------------|
| 5521 | 3863-QSTYI | Male   | 0             | No      | No         | 59     | Yes          |            |
| 480  | 0486-LGCCH | Male   | 0             | Yes     | Yes        | 11     | Yes          |            |
| 3929 | 9702-AIUJO | Male   | 0             | Yes     | Yes        | 50     | Yes          |            |
| 810  | 2239-JALAW | Male   | 0             | No      | No         | 58     | Yes          |            |
| 1822 | 4911-BANWH | Female | 0             | No      | Yes        | 31     | Yes          |            |

5 rows × 21 columns



In [3]:

```
df.drop('customerID',axis='columns',inplace=True)
```

In [4]:

```
df.dtypes
```

Out[4]:

```
gender          object
SeniorCitizen   int64
Partner         object
Dependents      object
tenure          int64
PhoneService    object
MultipleLines   object
InternetService object
OnlineSecurity  object
OnlineBackup    object
DeviceProtection object
TechSupport     object
StreamingTV     object
StreamingMovies object
Contract        object
PaperlessBilling object
PaymentMethod   object
MonthlyCharges  float64
TotalCharges    object
Churn           object
dtype: object
```

In [5]:

```
df.TotalCharges
```

Out[5]:

```
0      29.85
1     1889.5
2      108.15
3     1840.75
4      151.65
...
7038   1990.5
7039   7362.9
7040    346.45
7041     306.6
7042   6844.5
Name: TotalCharges, Length: 7043, dtype: object
```

In [6]:

```
df.TotalCharges.values
```

Out[6]:

```
array(['29.85', '1889.5', '108.15', ..., '346.45', '306.6', '6844.5'],
      dtype=object)
```

In [7]:

```
pd.to_numeric(df.TotalCharges,errors='coerce').isnull().sum()
```

Out[7]:

11

In [8]:

```
df.iloc[488].TotalCharges
```

Out[8]:

11

In [9]:

```
df1 = df[df.TotalCharges!=' ']  
df1.shape
```

Out[9]:

(7032, 20)

In [55]:

```
df1 = df[df.TotalCharges!=' ']  
df1.shape
```

Out[55]:

(7032, 20)

In [56]:

```
df1.dtypes  
df1.TotalCharges = pd.to_numeric(df1.TotalCharges)
```

C:\Users\Lenovo\anaconda3\lib\site-packages\pandas\core\generic.py:5516: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
self[name] = value
```

In [33]:

```
df1.dtypes
```

Out[33]:

```
gender          object
SeniorCitizen   int64
Partner         object
Dependents      object
tenure          int64
PhoneService    object
MultipleLines   object
InternetService object
OnlineSecurity  object
OnlineBackup    object
DeviceProtection object
TechSupport     object
StreamingTV     object
StreamingMovies object
Contract        object
PaperlessBilling object
PaymentMethod   object
MonthlyCharges  float64
TotalCharges    float64
Churn           object
dtype: object
```

In [57]:

```
df1.TotalCharges.values
```

Out[57]:

```
array([ 29.85, 1889.5 , 108.15, ..., 346.45, 306.6 , 6844.5 ])
```


In [58]:

```
df1[df1.Churn=='No']
```

Out[58]:

|      | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines    | Internet |
|------|--------|---------------|---------|------------|--------|--------------|------------------|----------|
| 0    | Female | 0             | Yes     | No         | 1      | No           | No phone service |          |
| 1    | Male   | 0             | No      | No         | 34     | Yes          | No               |          |
| 3    | Male   | 0             | No      | No         | 45     | No           | No phone service |          |
| 6    | Male   | 0             | No      | Yes        | 22     | Yes          | Yes              | Fit      |
| 7    | Female | 0             | No      | No         | 10     | No           | No phone service |          |
| ...  | ...    | ...           | ...     | ...        | ...    | ...          | ...              |          |
| 7037 | Female | 0             | No      | No         | 72     | Yes          | No               |          |
| 7038 | Male   | 0             | Yes     | Yes        | 24     | Yes          | Yes              |          |
| 7039 | Female | 0             | Yes     | Yes        | 72     | Yes          | Yes              | Fit      |
| 7040 | Female | 0             | Yes     | Yes        | 11     | No           | No phone service |          |
| 7042 | Male   | 0             | No      | No         | 66     | Yes          | No               | Fit      |

5163 rows × 20 columns




In [59]:

```
df1[df1.Churn=='Yes']
```

Out[59]:

|      | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | Internet |
|------|--------|---------------|---------|------------|--------|--------------|---------------|----------|
| 2    | Male   | 0             | No      | No         | 2      | Yes          | No            |          |
| 4    | Female | 0             | No      | No         | 2      | Yes          | No            | Fit      |
| 5    | Female | 0             | No      | No         | 8      | Yes          | Yes           | Fit      |
| 8    | Female | 0             | Yes     | No         | 28     | Yes          | Yes           | Fit      |
| 13   | Male   | 0             | No      | No         | 49     | Yes          | Yes           | Fit      |
| ...  | ...    | ...           | ...     | ...        | ...    | ...          | ...           |          |
| 7021 | Male   | 0             | No      | No         | 12     | Yes          | No            |          |
| 7026 | Female | 0             | No      | No         | 9      | Yes          | No            |          |
| 7032 | Male   | 1             | No      | No         | 1      | Yes          | Yes           | Fit      |
| 7034 | Female | 0             | No      | No         | 67     | Yes          | Yes           | Fit      |
| 7041 | Male   | 1             | Yes     | No         | 4      | Yes          | Yes           | Fit      |

1869 rows × 20 columns



In [60]:


```
tenure_churn_no = df1[df1.Churn=='No'].tenure
tenure_churn_yes = df1[df1.Churn=='Yes'].tenure

plt.xlabel("tenure")
plt.ylabel("Number Of Customers")
plt.title("Customer Churn Prediction Visualiztion")

plt.hist([tenure_churn_yes, tenure_churn_no], rwidth=0.95, color=['green', 'red'], label=['Ch
plt.legend()
```

Out[60]:

<matplotlib.legend.Legend at 0x274dee11160>



In [61]:


```
mc_churn_no = df1[df1.Churn=='No'].MonthlyCharges
mc_churn_yes = df1[df1.Churn=='Yes'].MonthlyCharges

plt.xlabel("Monthly Charges")
plt.ylabel("Number Of Customers")
plt.title("Customer Churn Prediction Visualiztion")

plt.hist([mc_churn_yes, mc_churn_no], rwidth=0.95, color=['green', 'red'], label=['Churn=Yes', 'Churn=No'])
plt.legend()
```

Out[61]:

<matplotlib.legend.Legend at 0x274c5aae130>



In [62]:

```
def print_unique_col_values(df):
    for column in df:
        if df[column].dtypes=='object':
            print(f'{column}: {df[column].unique()}')
```



In [63]:

```
print_unique_col_values(df)
```

```
gender: ['Female' 'Male']
Partner: ['Yes' 'No']
Dependents: ['No' 'Yes']
PhoneService: ['No' 'Yes']
MultipleLines: ['No phone service' 'No' 'Yes']
InternetService: ['DSL' 'Fiber optic' 'No']
OnlineSecurity: ['No' 'Yes' 'No internet service']
OnlineBackup: ['Yes' 'No' 'No internet service']
DeviceProtection: ['No' 'Yes' 'No internet service']
TechSupport: ['No' 'Yes' 'No internet service']
StreamingTV: ['No' 'Yes' 'No internet service']
StreamingMovies: ['No' 'Yes' 'No internet service']
Contract: ['Month-to-month' 'One year' 'Two year']
PaperlessBilling: ['Yes' 'No']
PaymentMethod: ['Electronic check' 'Mailed check' 'Bank transfer (automatic)'
'Credit card (automatic)']
TotalCharges: ['29.85' '1889.5' '108.15' ... '346.45' '306.6' '6844.5']
Churn: ['No' 'Yes']
```

In [64]:

```
df1.replace('No internet service', 'No', inplace=True)
df1.replace('No phone service', 'No', inplace=True)
```

C:\Users\Lenovo\anaconda3\lib\site-packages\pandas\core\frame.py:5238: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
    return super().replace(
```

In [65]:

```
print_unique_col_values(df1)
```

```
gender: ['Female' 'Male']
Partner: ['Yes' 'No']
Dependents: ['No' 'Yes']
PhoneService: ['No' 'Yes']
MultipleLines: ['No' 'Yes']
InternetService: ['DSL' 'Fiber optic' 'No']
OnlineSecurity: ['No' 'Yes']
OnlineBackup: ['Yes' 'No']
DeviceProtection: ['No' 'Yes']
TechSupport: ['No' 'Yes']
StreamingTV: ['No' 'Yes']
StreamingMovies: ['No' 'Yes']
Contract: ['Month-to-month' 'One year' 'Two year']
PaperlessBilling: ['Yes' 'No']
PaymentMethod: ['Electronic check' 'Mailed check' 'Bank transfer (automatic)'
'Credit card (automatic)']
Churn: ['No' 'Yes']
```

In [66]:

```
for col in df1:
    print(f'{col}: {df1[col].unique()}')
```

```
gender: ['Female' 'Male']
SeniorCitizen: [0 1]
Partner: ['Yes' 'No']
Dependents: ['No' 'Yes']
tenure: [ 1 34  2 45  8 22 10 28 62 13 16 58 49 25 69 52 71 21 12 30 47 72 1
7 27
 5 46 11 70 63 43 15 60 18 66  9  3 31 50 64 56  7 42 35 48 29 65 38 68
32 55 37 36 41  6  4 33 67 23 57 61 14 20 53 40 59 24 44 19 54 51 26 39]
PhoneService: ['No' 'Yes']
MultipleLines: ['No' 'Yes']
InternetService: ['DSL' 'Fiber optic' 'No']
OnlineSecurity: ['No' 'Yes']
OnlineBackup: ['Yes' 'No']
DeviceProtection: ['No' 'Yes']
TechSupport: ['No' 'Yes']
StreamingTV: ['No' 'Yes']
StreamingMovies: ['No' 'Yes']
Contract: ['Month-to-month' 'One year' 'Two year']
PaperlessBilling: ['Yes' 'No']
PaymentMethod: ['Electronic check' 'Mailed check' 'Bank transfer (automatic)'
'Credit card (automatic)']
MonthlyCharges: [29.85 56.95 53.85 ... 63.1  44.2  78.7 ]
TotalCharges: [ 29.85 1889.5  108.15 ... 346.45  306.6  6844.5 ]
Churn: ['No' 'Yes']
```

In [67]:

```
df1.gender.unique()
```

Out[67]:

```
array(['Female', 'Male'], dtype=object)
```

In [176]:

```
df1.Churn.unique()
```

Out[176]:

```
array(['No', 'Yes'], dtype=object)
```

In [177]:

```
df1['gender'].replace({'Female':1, 'Male':0}, inplace=True)
```

C:\Users\Lenovo\anaconda3\lib\site-packages\pandas\core\generic.py:6619: Set  
tingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
return self._update_inplace(result)
```

In [178]:

```
df1['Partner'].replace({'Yes':1, 'No':0}, inplace=True)
```

In [179]:

```
df1['Dependents'].replace({'Yes':1, 'No':0}, inplace=True)
```

In [180]:

```
df1['PhoneService'].replace({'Yes':1, 'No':0}, inplace=True)
```

In [181]:

```
df1['MultipleLines'].replace({'Yes':1, 'No':0}, inplace=True)
```

In [182]:

```
df1['OnlineSecurity'].replace({'Yes':1, 'No':0}, inplace=True)
```

In [183]:

```
df1['OnlineBackup'].replace({'Yes':1, 'No':0}, inplace=True)
```

In [184]:

```
df1['OnlineSecurity'].replace({'Yes':1, 'No':0}, inplace=True)
```

In [185]:

```
df1['DeviceProtection'].replace({'Yes':1, 'No':0}, inplace=True)
df1['Churn'].replace({'Yes':1, 'No':0}, inplace=True)
```

In [186]:

```
df1['TechSupport'].replace({'Yes':1, 'No':0}, inplace=True)
df1['StreamingTV'].replace({'Yes':1, 'No':0}, inplace=True)
df1['StreamingMovies'].replace({'Yes':1, 'No':0}, inplace=True)
df1['PaperlessBilling'].replace({'Yes':1, 'No':0}, inplace=True)
```

In [187]:

```
df2 = pd.get_dummies(data=df1, columns=['InternetService', 'Contract', 'PaymentMethod'])
df2.columns
```

Out[187]:

```
Index(['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure',
      'PhoneService', 'MultipleLines', 'OnlineSecurity', 'OnlineBackup',
      'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies',
      'PaperlessBilling', 'MonthlyCharges', 'TotalCharges', 'Churn',
      'InternetService_DSL', 'InternetService_Fiber optic',
      'InternetService_No', 'Contract_Month-to-month', 'Contract_One year',
      'Contract_Two year', 'PaymentMethod_Bank transfer (automatic)',
      'PaymentMethod_Credit card (automatic)',
      'PaymentMethod_Electronic check', 'PaymentMethod_Mailed check'],
      dtype='object')
```


In [188]:

```
df2.sample(5)
```

Out[188]:

|      | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | OnlineS |
|------|--------|---------------|---------|------------|--------|--------------|---------------|---------|
| 863  | 1      | 0             | 0       | 0          | 3      | 1            | 0             |         |
| 4104 | 0      | 0             | 1       | 0          | 8      | 1            | 0             |         |
| 6355 | 1      | 0             | 0       | 0          | 68     | 1            | 1             |         |
| 3678 | 0      | 0             | 1       | 1          | 16     | 1            | 1             |         |
| 1456 | 1      | 0             | 1       | 0          | 16     | 1            | 1             |         |

5 rows × 27 columns



In [189]:

```
df2.dtypes
```

Out[189]:

|   |         |
|---|---------|
| gender                                  | int64   |
| SeniorCitizen                           | int64   |
| Partner                                 | int64   |
| Dependents                              | int64   |
| tenure                                  | int64   |
| PhoneService                            | int64   |
| MultipleLines                           | int64   |
| OnlineSecurity                          | int64   |
| OnlineBackup                            | int64   |
| DeviceProtection                        | int64   |
| TechSupport                             | int64   |
| StreamingTV                             | int64   |
| StreamingMovies                         | int64   |
| PaperlessBilling                        | int64   |
| MonthlyCharges                          | float64 |
| TotalCharges                            | float64 |
| Churn                                   | int64   |
| InternetService_DSL                     | uint8   |
| InternetService_Fiber optic             | uint8   |
| InternetService_No                      | uint8   |
| Contract_Month-to-month                 | uint8   |
| Contract_One year                       | uint8   |
| Contract_Two year                       | uint8   |
| PaymentMethod_Bank transfer (automatic) | uint8   |
| PaymentMethod_Credit card (automatic)   | uint8   |
| PaymentMethod_Electronic check          | uint8   |
| PaymentMethod_Mailed check              | uint8   |
| dtype:                                  | object  |

In [190]:

```
cols_to_scale = ['tenure', 'MonthlyCharges', 'TotalCharges']

from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
df2[cols_to_scale] = scaler.fit_transform(df2[cols_to_scale])
```

In [191]:

```
for col in df2:
    print(f'{col}: {df2[col].unique()}')
```

```
gender: [1 0]
SeniorCitizen: [0 1]
Partner: [1 0]
Dependents: [0 1]
tenure: [0.         0.46478873 0.01408451 0.61971831 0.09859155 0.29577465
 0.12676056 0.38028169 0.85915493 0.16901408 0.21126761 0.8028169
 0.67605634 0.33802817 0.95774648 0.71830986 0.98591549 0.28169014
 0.15492958 0.4084507  0.64788732 1.         0.22535211 0.36619718
 0.05633803 0.63380282 0.14084507 0.97183099 0.87323944 0.5915493
 0.1971831  0.83098592 0.23943662 0.91549296 0.11267606 0.02816901
 0.42253521 0.69014085 0.88732394 0.77464789 0.08450704 0.57746479
 0.47887324 0.66197183 0.3943662  0.90140845 0.52112676 0.94366197
 0.43661972 0.76056338 0.50704225 0.49295775 0.56338028 0.07042254
 0.04225352 0.45070423 0.92957746 0.30985915 0.78873239 0.84507042
 0.18309859 0.26760563 0.73239437 0.54929577 0.81690141 0.32394366
 0.6056338  0.25352113 0.74647887 0.70422535 0.35211268 0.53521127]
PhoneService: [0 1]
MultipleLines: [0 1]
OnlineSecurity: [0 1]
OnlineBackup: [1 0]
DeviceProtection: [0 1]
TechSupport: [0 1]
StreamingTV: [0 1]
StreamingMovies: [0 1]
PaperlessBilling: [1 0]
MonthlyCharges: [0.11542289 0.38507463 0.35422886 ... 0.44626866 0.25820896
0.60149254]
TotalCharges: [0.0012751 0.21586661 0.01031041 ... 0.03780868 0.03321025 0.
78764136]
Churn: [0 1]
InternetService_DSL: [1 0]
InternetService_Fiber optic: [0 1]
InternetService_No: [0 1]
Contract_Month-to-month: [1 0]
Contract_One year: [0 1]
Contract_Two year: [0 1]
PaymentMethod_Bank transfer (automatic): [0 1]
PaymentMethod_Credit card (automatic): [0 1]
PaymentMethod_Electronic check: [1 0]
PaymentMethod_Mailed check: [0 1]
```

In [192]:

```
X = df2.drop('Churn',axis='columns')
y = df2['Churn']

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2,random_state=5)
```

In [193]:

```
X_train.shape
```

Out[193]:

```
(5625, 26)
```

In [194]:

```
X_test.shape
```

Out[194]:

```
(1407, 26)
```

In [195]:

```
X_train = X_train.astype(float)
```

In [ ]:

In [ ]:

In [ ]:

In [196]:

```
import tensorflow as tf
from tensorflow import keras

model = keras.Sequential([
    keras.layers.Dense(26, input_shape=(26,), activation='relu'),
    keras.layers.Dense(15, activation='relu'),
    keras.layers.Dense(1, activation='sigmoid')
])

model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['accuracy'])

model.fit(X_train, y_train, epochs=100)

176/176 [=====] - 0s 2ms/step - loss: 0.3508 - a
ccuracy: 0.8292
Epoch 85/100
176/176 [=====] - 0s 2ms/step - loss: 0.3592 - a
ccuracy: 0.8304
Epoch 86/100
176/176 [=====] - 0s 2ms/step - loss: 0.3582 - a
ccuracy: 0.8299
Epoch 87/100
176/176 [=====] - 0s 2ms/step - loss: 0.3583 - a
ccuracy: 0.8357
Epoch 88/100
176/176 [=====] - 0s 2ms/step - loss: 0.3575 - a
ccuracy: 0.8334
Epoch 89/100
176/176 [=====] - 0s 2ms/step - loss: 0.3570 - a
ccuracy: 0.8320
Epoch 90/100
176/176 [=====] - 0s 2ms/step - loss: 0.3548 - a
ccuracy: 0.8345
- . . . . .
```

In [197]:

```
model.evaluate(X_test, y_test)
```

```
44/44 [=====] - 0s 1ms/step - loss: 0.4905 - accura
cy: 0.7726
```

Out[197]:

```
[0.49049249291419983, 0.7725657224655151]
```



In [200]:

```
yp = model.predict(X_test)
yp[:5]
```

44/44 [=====] - 0s 1ms/step

Out[200]:

```
array([[0.28224123],
       [0.4679681 ],
       [0.00627228],
       [0.8261105 ],
       [0.5812232 ]], dtype=float32)
```

In [201]:

```
y_pred = []
for element in yp:
    if element > 0.5:
        y_pred.append(1)
    else:
        y_pred.append(0)
y_pred[:10]
[0, 0, 0, 1, 0, 1, 0, 0, 0, 0]
y_test[:10]
```

Out[201]:

```
2660    0
744     0
5579    1
64      1
3287    1
816     1
2670    0
5920    0
1023    0
6087    0
Name: Churn, dtype: int64
```

In [202]:

```
from sklearn.metrics import confusion_matrix , classification_report
print(classification_report(y_test,y_pred))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.81      | 0.88   | 0.85     | 999     |
| 1            | 0.64      | 0.50   | 0.56     | 408     |
| accuracy     |           |        | 0.77     | 1407    |
| macro avg    | 0.73      | 0.69   | 0.70     | 1407    |
| weighted avg | 0.76      | 0.77   | 0.76     | 1407    |


In [203]:

```
import seaborn as sn
cm = tf.math.confusion_matrix(labels=y_test,predictions=y_pred)

plt.figure(figsize = (10,7))
sn.heatmap(cm, annot=True, fmt='d')
plt.xlabel('Predicted')
plt.ylabel('Truth')
```

Out[203]:

Text(69.0, 0.5, 'Truth')



In [ ]: